

健康資訊交換系統中之容器安全
The Container Security in Healthcare Data Exchange System

國立中山大學資訊工程學系
Department of Computer Science and Engineering
National Sun-Yet-San University, Taiwan
110 學年度大學部專題製作競賽
Bachelor's degree graduation project in 2021

Author: Chih-Hsuan Yang (B073040047)

Advisor: Chun-I Fan

October 12, 2021

Abstract

This research proposes a mechanism to enforce the system call a specific policy in the container, which is deployed in runtime. This policy is designed for the FHIR healthcare data exchange standard's container, which could guarantee the FHIR server does not have unsupported behavior and takes almost zero overhead. Recently, many companies use containers to run their microservices since containers could make their hardware resources be used efficiently. And the newest healthcare data exchange standard FHIR (Fast Healthcare Interoperability Resources) ¹ has been implemented in a container by IBM, Microsoft, and Firebase. The deployment of FHIR in a container is a trend in the digital world [1]. However, containers are not sandboxes [2]. Containers are just isolated processes ². Therefore, if hackers or malicious software could sneak into the container that would be a new cyber attacking surface in nearly future.

¹FHIR official: <https://www.hl7.org/fhir/>

²gVisor GitHub: <https://github.com/google/gvisor>

Contents

Abstract	1
1 Introduction	4
1.1 Container and Linux Kernel	4
1.2 FHIR	5
1.2.1 RESTful API and Data Structure	5
1.2.2 Why IBM FHIR server	5
1.3 Data and Privacy	5
2 Related Work	6
2.1 Collecting System Calls	6
2.2 Fine-grained Permission Control	7
2.2.1 Capabilities	7
2.2.2 Linux Secure Module	7
2.3 Recently Exploited Vulnerabilities	8
2.3.1 Five Stages of Malware	8
2.4 Virtual Environment Performance Benchmark	8
3 Preliminary	9
3.1 Container's Components	9
3.1.1 Namespaces	9
3.1.2 Cgroups	9
3.1.3 Seccomp	9
3.2 Programs in Execution	9
3.2.1 The task_struct in Kernel	9
3.2.2 Capabilities	9
3.3 Sandbox Security	9
3.3.1 User Mode Linux	9
3.3.2 Virtual Machines	9
3.4 The (e)BPF	9
4 Proposed Scheme	10
4.1 Workflow	10
4.1.1 Scan Base Image	10
4.1.2 Building and Signing	10
4.1.3 Check Image and Policy	10
4.1.4 Enforce the Policy	10
4.2 Rolling Updates	10
5 Analysis and Benchmark	11
5.1 Analysis	11
5.1.1 Attacking Surface	11
5.1.2 Time Consuming	11
5.1.3 Statistics	11
5.2 Benchmark	11
5.2.1 Latency	11
5.2.2 Throughput	11

6 Conclusion **12**

6.1 Better Architecture 12

6.2 Future Machine Learning in Kernel 12

Reference **13**

Chapter 1

Introduction

1.1 Container and Linux Kernel

The container is a secondary product of the operating system in the past 20 years. The FreeBSD develops ‘Jails’ in 1999, and the Solaris develops ‘Zones’ in 2004. Linux also took this idea into the Linux kernel, which is named cgroups (2007), the capabilities (2003), and seccomp (2005). However, why the Linux breaks this technology into many parts? This is because they had discussed: “Why Should a System Administrator Upgrade?” in 2001 ¹. The Linux kernel almost entered the development path of “upgrade for demand” like Microsoft Windows, and deviated from the original path of “providing a mechanism but not a strategy” of the original Linux kernel.

While Linux were spreading in various server or distributed system, the Linux community got more pull requests to solved the scalability and virtualization issues [3]. However, they avoided confusion caused by multiple meanings of the term “container” in the Linux kernel context. In kernel version 2.6.24 (2007) ², control groups functionality was merged into the mainline, which is designed for an administrator (or administrative daemon) to organize processes into hierarchies of containers; each hierarchy is managed by a subsystem. Moreover, the cgroups was rewrote into cgroups-v2 in Linux kernel 4.5 (2015) ³.

The first and most complete implementation of the Linux container manager was LXC (Linux Containers). It was implemented in 2008 using cgroups and namespaces, and it runs on a single Linux kernel without requiring any patches. LXC provides a new view and imagination of virtualized services without any hypervisor. In 2016, Docker replaced LXC with “libcontainer”, which was written in the Go programming language. Docker combined features in a new, more attractive way and made Linux containers popular.

The secondary product of the operating system, containers, offering many advantages: they enable you to “build once, run anywhere.” Docker does this by bundling applications with all their dependencies into one package and isolating applications from the rest of the machine on which they’re running. Therefore, this research is based on docker container to propose a scheme of healthcare data exchange system’s security.

¹Version 2.4 of the LINUX KERNEL–Why Should a System Administrator Upgrade? <https://www.informit.com/articles/article.aspx?p=20667>

²Notes from a container: <https://lwn.net/Articles/256389/>

³Control Group v2: <https://www.kernel.org/doc/Documentation/cgroup-v2.txt>

1.2 FHIR

FHIR is a standard for healthcare data exchange. The FHIR standard will be used in Taiwan in the near future. FHIR will be used to provide PHR (Personal Healthcare Records) in Taiwan. Therefore, we choose the most popular standard "FHIR" for the target of the healthcare data exchange system.

1.2.1 RESTful API and Data Structure

REST (Representational State Transfer) is a stateless reliable web API, which is based on HTTP methods to access resources or data via URL parameters and the use of JSON or XML format to transmit queries. Because the RESTful is stateless, the client should keep their information (i.e. cookies) by themselves.

FHIR has features: RESTful and data structure, make our research and benchmarks more accurate and reliable. Statelessness is a developer-friendly feature, the developer and the tester would not to design a complex state machine on the server-side or generating test files. And the FHIR takes RESTful as standard. Moreover, FHIR standard declared the 'StructureDefinition' ⁴. These structure definitions are used to describe both the content defined in the FHIR specification itself - Resources, data types, the underlying infrastructural types, and also are used to describe how these structures are used in implementations.

1.2.2 Why IBM FHIR server

There are many applications using IBM's FHIR server as the base component of the EHR (Electronic Health Records) system to communicate with the other various databases. Take it for example that the NextCloud's EHR service, Taipei Veterans General Hospital, and AWS Cloud are using the FHIR server in a container for subroutine service.

NextCloud is an open-source and self-hosted productivity platform for users. Many people caring about their privacy issues distrust the FAANG (Facebook, Amazon, Apple, Netflix, Google), so they are using NextCloud to keep their privacy on their own. Therefore, they are eager to have a secure EHR system for their PHR ⁵.

1.3 Data and Privacy

TBD.

⁴FHIR Resource Structure Definition: <http://www.hl7.org/fhir/structuredefinition.html>

⁵Richard Stallman talks about IoT

Chapter 2

Related Work

2.1 Collecting System Calls

There are several pieces of research to detect intrusions or unexpected behaviors by collecting the system calls methods in runtime [4, 5, 6, 7]. Amr S. Abed et al.[4] proposed a real-time host-based intrusion detection system in a container, which is based on system call monitoring. They use the ‘strace’ command to collect a behavior log to a system call parser. Then use the BoSC (Bag of System Calls) [8] to classify is it a normal behavior in the database.

The BoSC technique is a frequency-based detection tip. Kang et al.[8] defined those distinct system calls in $\{c_1, c_2, \dots, c_n\}$, \forall system call s_i had been called in c_i times. And they use Naïve Bayes classification to deduce if it is unexpected behavior. Then the Amr S. Abed et al. give the false positive rate around 2% in $O(S + n_k)$ epochs to the MySQL database [4].

- Epoch Size (S): The total number of system calls in one epoch.
- n_k : It is the size of the database after epoch k .

However, the BoSC is running in user space, even though it is a background service running on the same host kernel. It might have heavy constant time costs of copying data from user to kernel and kernel to user by the ‘copy_to_user()’ and ‘copy_from_user()’ calls.

Mohamed Azab et al.[6, 7] takes a mathematical model to simulate the smart moving target defense for Linux container resiliency. Considering an ‘ESCAPE’ model is the interaction between attackers and their target containers as a “predator searching for a prey” search game. This search game has 3 modules: behavior monitoring, the checkpoint/restore, and the live migration modules. This model is running on the same host and the same attacking surface because they considered the containers (prey) are running on the same machine with some migration probability.

They show the survival rate in Amr S. Abed et al.[4] model for some zero-day vulnerabilities in different types and numbers machines. Mohamed Azab et al. [6, 7] concluded that an IDS could detect and avoid mobile continually-growing attacks efficiently by the ‘ESCAPE’ model with collecting system calls.

2.2 Fine-grained Permission Control

2.2.1 Capabilities

There are 49 different capabilities in today's Linux kernel 5.13¹. A capability can be assigned to a task (i.e. thread or process) to determine if the task can use the fine-granted system calls. For example, we give a thread `CAP_SYS_BOOT`, then the thread can use the `reboot`² and the `kexec_load`³ system call.

Xin Lin et al.[9] collected 27 CVE vulnerabilities could cause the privilege escalation attacks. There are only 3 vulnerabilities could bypass the capabilities protection of Linux kernel. And the other 24 escalation vulnerabilities, could be filtered by the fine-grained permission control with capabilities. Those 3 (CVE-2016-8655, CVE-2017-5123, CVE-2017-7308) bypassed capabilities vulnerabilities are attacking kernel level race conditions.

The CVE-2016-8655⁴ is a bug in `net/packet/af_packet.c`. We often use the `CAP_NET_RAW` namespace in the container to make unprivileged user be able to use some privileged net-util commands. The bug is that there exists a race condition probability to race the unauthorized data inside `packet_set_ring()` and `packet_setsockopt()`⁵ such that there is a chance to modify the socket version to `TPACKET_V1` before the `packet_set_ring` function. However, it would be 'kfree' the timer in the `TPACKET_V1`. Then we can take the timer, which is used after free, to control the SLAB allocator to write the `st_uid` by itself⁶.

2.2.2 Linux Secure Module

There were many researches designing great rules of the LSM (Linux secure module). LSM is a mandatory access control framework in kernel, which supported from Linux 2.6 (2004). The AppArmor⁷ and SELinux⁸ are common LSM in the kernel. All the Android devices are fully enforced the SELinux after Android 5.0.⁹

The SIDs (Security IDs) and permissions, which are the identifiers for access control policies, make up the security policy of SELinux or AppArmor. Files and directories, which are the actual protected objects for the SID, are mapped to the SELinux or AppArmor of each system [10, 11, 12]. The SELinux is much more incredibly complex than AppArmor, but with this complexity you have more control over how tasks are isolated. However, the AppArmor is so straightforward that people can write the configuration profile by themselves.

SUNG-HWA HAN et al.[13] had proposed an architecture to enforce the access control of image's layers. Because the docker engine does not guarantee the layers could not be modified by the host environment. Therefore, if we give a container privileged permission, it could modify the layers of images. The research [13] is using the LSM's policy table to enforce the access control of file system in kernel.

This paper [13] shows that the performance is almost same as raw SELinux, and the branch costs of indexing the policy table could be a constant value in the CPU rate measurement. However, the research is only measured in over-

¹<https://man7.org/linux/man-pages/man7/capabilities.7.html>

²<https://man7.org/linux/man-pages/man2/reboot.2.html>

³https://man7.org/linux/man-pages/man2/kexec_load.2.html

⁴<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2016-8655>

⁵Linux `af_packet.c` race condition (local root) <https://lwn.net/Articles/708319/>

⁶CVE-2016-8655 PoC: <https://www.exploit-db.com/exploits/40871>

⁷<https://apparmor.net/>

⁸<https://github.com/SELinuxProject/selinux>

⁹https://source.android.com/security/selinux#supporting_documentation

lay2 file system. SUNG-HWA HAN et al. said if it have same performance in AUFS or the other file systems, the above results could be more reliable.

Yuqiong Sun et al.[14] proposed separate the security namespace. Each container can route their operation to different security namespaces for their "comment". Each involved the security namespace independently makes a security decision, and the operation is allowed only if all the secure namespaces allow.

The policy engine and the security namespace are implemented in ...

2.3 Recently Exploited Vulnerabilities

2.3.1 Five Stages of Malware

2.4 Virtual Environment Performance Benchmark

Chapter 3

Preliminary

3.1 Container's Components

3.1.1 Namespaces

3.1.2 Cgroups

3.1.3 Seccomp

3.2 Programs in Execution

3.2.1 The `task_struct` in Kernel

3.2.2 Capabilities

3.3 Sandbox Security

3.3.1 User Mode Linux

gVisor

3.3.2 Virtual Machines

3.4 The (e)BPF

Chapter 4

Proposed Scheme

4.1 Workflow

4.1.1 Scan Base Image

4.1.2 Building and Signing

4.1.3 Check Image and Policy

4.1.4 Enforce the Policy

4.2 Rolling Updates

Chapter 5

Analysis and Benchmark

5.1 Analysis

5.1.1 Attacking Surface

5.1.2 Time Consuming

5.1.3 Statistics

5.2 Benchmark

5.2.1 Latency

5.2.2 Throughput

Chapter 6

Conclusion

6.1 Better Architecture

6.2 Future Machine Learning in Kernel

Reference

- [1] Arif Ahmed and Guillaume Pierre. “Docker Container Deployment in Fog Computing Infrastructures”. In: *2018 IEEE International Conference on Edge Computing (EDGE)*. 2018, pp. 1–8. DOI: [10.1109/EDGE.2018.00008](https://doi.org/10.1109/EDGE.2018.00008).
- [2] Ian Goldberg et al. “A Secure Environment for Untrusted Helper Applications Confining the Wily Hacker”. In: *Proceedings of the 6th Conference on USENIX Security Symposium, Focusing on Applications of Cryptography - Volume 6*. SSYM’96. San Jose, California: USENIX Association, 1996, p. 1.
- [3] Silas Boyd-Wickizer et al. “An Analysis of Linux Scalability to Many Cores”. In: *9th USENIX Symposium on Operating Systems Design and Implementation (OSDI 10)*. Vancouver, BC: USENIX Association, Oct. 2010. URL: <https://www.usenix.org/conference/osdi10/analysis-linux-scalability-many-cores>.
- [4] Amr Abed, Charles Clancy, and David Levy. “Intrusion Detection System for Applications Using Linux Containers”. In: vol. 9331. Sept. 2015, pp. 123–135. ISBN: 978-3-319-24857-8. DOI: [10.1007/978-3-319-24857-8_5](https://doi.org/10.1007/978-3-319-24857-8_5).
- [5] José Flora. “Improving the Security of Microservice Systems by Detecting and Tolerating Intrusions”. In: *2020 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*. 2020, pp. 131–134. DOI: [10.1109/ISSREW51248.2020.00051](https://doi.org/10.1109/ISSREW51248.2020.00051).
- [6] Mohamed Azab et al. “Smart Moving Target Defense for Linux Container Resiliency”. In: *2016 IEEE 2nd International Conference on Collaboration and Internet Computing (CIC)*. 2016, pp. 122–130. DOI: [10.1109/CIC.2016.028](https://doi.org/10.1109/CIC.2016.028).
- [7] Mohamed Azab et al. “Toward Smart Moving Target Defense for Linux Container Resiliency”. In: *2016 IEEE 41st Conference on Local Computer Networks (LCN)*. 2016, pp. 619–622. DOI: [10.1109/LCN.2016.106](https://doi.org/10.1109/LCN.2016.106).
- [8] Dae-Ki Kang, D. Fuller, and V. Honavar. “Learning classifiers for misuse and anomaly detection using a bag of system calls representation”. In: *Proceedings from the Sixth Annual IEEE SMC Information Assurance Workshop*. 2005, pp. 118–125. DOI: [10.1109/IAW.2005.1495942](https://doi.org/10.1109/IAW.2005.1495942).
- [9] Xin Lin et al. “A Measurement Study on Linux Container Security: Attacks and Countermeasures”. In: *Proceedings of the 34th Annual Computer Security Applications Conference*. ACSAC ’18. San Juan, PR, USA: Association for Computing Machinery, 2018, pp. 418–429. ISBN: 9781450365697. DOI: [10.1145/3274694.3274720](https://doi.org/10.1145/3274694.3274720). URL: <https://doi.org/10.1145/3274694.3274720>.
- [10] Stephen Dale Smalley, Chris Vance, and Wayne Salamon. “Implementing SELinux as a Linux Security Module”. In: 2003.
- [11] Doug Kilpatrick, Wayne Salamon, and Chris Vance. “Securing The X Window System With SELinux”. In: 2003.
- [12] Luis Franco, Tony Sahama, and Peter Croll. “Security Enhanced Linux to enforce Mandatory Access Control in Health Information Systems”. In: *Health Data and Knowledge Management 2008*. Ed. by P Yu P et al. Australia: Australian Computer Society, 2008, pp. 27–33. URL: <https://eprints.qut.edu.au/30563/>.
- [13] Sung-Hwa Han et al. “Container Image Access Control Architecture to Protect Applications”. In: *IEEE Access* 8 (2020), pp. 162012–162021. DOI: [10.1109/ACCESS.2020.3021044](https://doi.org/10.1109/ACCESS.2020.3021044).
- [14] Yuqiong Sun et al. “Security Namespace: Making Linux Security Frameworks Available to Containers”. In: *27th USENIX Security Symposium (USENIX Security 18)*. Baltimore, MD: USENIX Association, Aug. 2018, pp. 1423–1439. ISBN: 978-1-939133-04-5. URL: <https://www.usenix.org/conference/usenixsecurity18/presentation/sun>.
- [15] Xing Gao et al. “ContainerLeaks: Emerging Security Threats of Information Leakages in Container Clouds”. In: *2017 47th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*. 2017, pp. 237–248. DOI: [10.1109/DSN.2017.49](https://doi.org/10.1109/DSN.2017.49).