

Docker

Information security club course II

Chih-Hsuan Yang(SCC)

National Sun Yat-sen University

March 5, 2021

v0.1

Outline

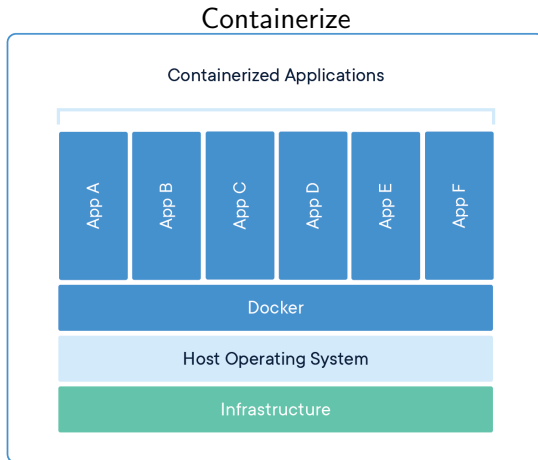
- 1 Container
 - Basic
 - Comparing Containers and Virtual Machines
- 2 Dockerfile
 - Introduction
 - Write a Dockerfile
- 3 Docker commands
 - Basic commands
 - Exercises
- 4 Tiny project
 - Image deployment
 - X Window forwarding
 - Maybe: Cross platform
- 5 Security
- 6 References

Before the speech

- The 'Learning Corner TA' of operating system course
 - 18:00 → 21:00 (Thur.)
 - EC1013
 - From 25nd March 2021
- The other time for OS problems
 - 18:00 → 21:00 (Thur.)
 - EC3034
 - Sent Email before you come.
- zxc25077667@protonmail.com
- We learn more only if you ask.

Container

Big idea

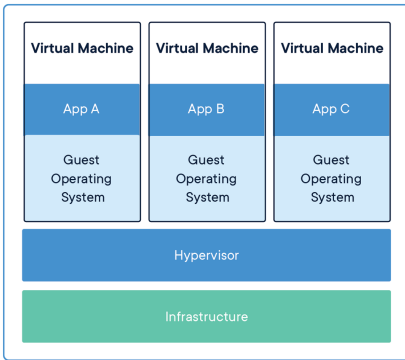
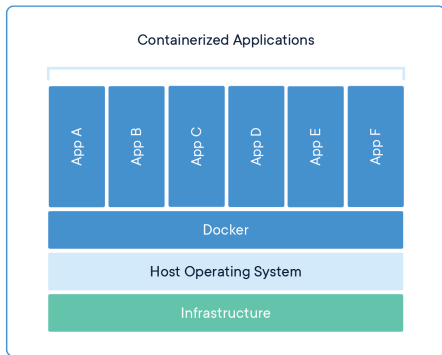


[1]

Big idea




Compare with virtule machines



So, what is share kernel?

Let's recall the OS 101 course

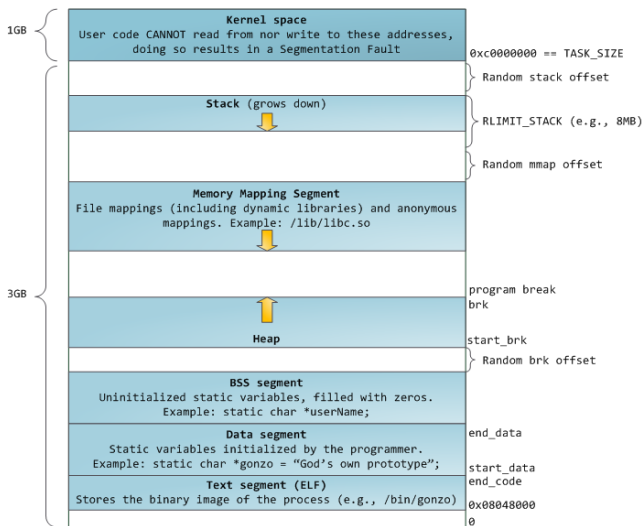


So Kernel + User Spaces add for 256 TiB which is a tiny part of the 16 777 216 TiB addressable over 64 bit!

[3]

So, what is share kernel?

The 32-bits memory layout



[4]

What is HYPERVISOR?

Virtual machine monitor


Dockerfile

What is Dockerfile?

Definition:

A **text** document that contains all the **commands** a user could call on the command line to **assemble an image**.

Image and Container



[5]

Like a program in execution is called a process.
An image in execution is called a container.

DockerHub

A public **images** hub. We push/pull **images** from it by default.
We push an image and pull an image rather than a container.

Make sense, right?

Dockerfile 101

Our Dockerfile

Docker commands

Basic



Commands Cheat Sheet

Container Lifecycle

docker create [IMAGE]	create a container without starting it
docker rename [CONTAINER_NAME] [NEW_CONTAINER_NAME]	rename a container
docker run [IMAGE]	create and start a container
docker run --rm [IMAGE]	remove a container after it stops
docker run -td [IMAGE]	start a container and keep it running
docker run -it [IMAGE]	create, start the container, and run a command in it
docker run -it-rm [IMAGE]	create, start the container, and run a command in it; after executing, the container is removed
docker rm [CONTAINER]	delete a container if it isn't running
docker update [CONTAINER]	update the configuration of a container

Networking

docker network ls	list networks
docker network rm [NETWORK]	remove one or more networks
docker network inspect [NETWORK]	show information on one or more networks
docker network connect [NETWORK] [CONTAINER]	connect a container to a network
docker network disconnect [NETWORK] [CONTAINER]	disconnect a container from a network

Image Lifecycle

docker build [URL]	create an image from a Dockerfile
docker build -t [URL]	build an image from a Dockerfile and tags it
docker pull [IMAGE]	pull an image from a registry
docker push [IMAGE]	push an image to a registry
docker import [URL/FILE]	create an image from a tarball
docker commit [CONTAINER] [NEW_IMAGE_NAME]	create an image from a container
docker rmi [IMAGE]	remove an image
docker load [TAR_FILE/STDIN_FILE]	load an image from a tar archive as stdin
docker save [IMAGE] > [TAR_FILE]	save an image to a tar archive stream to stdout with all parent layers, tags, and versions

Start & Stop

docker start [CONTAINER]	start a container
docker stop [CONTAINER]	stop a running container
docker restart [CONTAINER]	stop a running container and start it up again
docker pause [CONTAINER]	pause processes in a running container
docker unpause [CONTAINER]	unpause processes in a container
docker wait [CONTAINER]	block a container until other containers stop
docker kill [CONTAINER]	kill a container by sending SIGKILL to a running container
docker attach [CONTAINER]	attach local standard input, output, and error streams to a running container

Information

docker ps	list running containers
docker ps -a	list running and stopped containers
docker logs [CONTAINER]	list the logs from a running container
docker inspect [OBJECT_NAME/ID]	list low-level information on an object
docker events [CONTAINER]	list real time events from a container
docker port [CONTAINER]	show port (or specific) mapping from a container
docker top [CONTAINER]	show running processes in a container
docker stats [CONTAINER]	show live resource usage statistics of containers
docker diff [CONTAINER]	show changes to files (or directories) on a filesystem
docker images ls	show all locally stored images
docker history [IMAGE]	show history of an image



Full cheat sheet

Next page reference: [7]

1. Containers

A lightweight virtual OS that runs processes in full isolation.

1.1 Lifecycle

- `docker create` creates a container but does not start it.
- `docker rename` allows the container to be renamed.
- `docker run` creates and starts a container in one operation.
- `docker rm` deletes a container.
- `docker update` updates a container's resource limits.
- `docker run --rm`: remove the container after it stops.
- `docker run -v SHOSTDIR:SDOCKERDIR`: map the directory (SHOSTDIR) on the host to a docker container (SDOCKERDIR).
- `docker rm --v`: remove the volumes associated with the container.
- `docker run --log-driver=syslog`: run docker with a custom log driver.

1.2 Starting and Stopping

- `docker start` starts a container so it is running.
- `docker stop` stops a running container.
- `docker restart` stops and starts a container.
- `docker pause` pauses a running container, "freezing" it in place.
- `docker unpause` will unpause a running container.
- `docker wait` blocks until running container stops.
- `docker kill` sends a SIGKILL to a running container.
- `docker attach` will connect to a running container.

1.3 CPU Constraints

CPU can be limited either using a percentage over all CPUs, or by using specific cores.

- `-c` or `-cpu-shares`: 1024 means 100% of the CPU, so if we want the container to take 50% of all CPU cores, we should specify 512 for instance, `docker run -t -c 512 ...cpuset-cpus`
- use only some CPU cores, for instance, `docker run -t --cpuset-cpus=0,4,6 ...`

1.4 Memory Constraints

Memory can be limited using `-m` flag, for instance, `docker run -it -m 300M ubuntu:14.04 /bin/bash`

1.5 Capabilities

`cap-add` and `cap-drop`: Add or drop linux capabilities.

- Mount a FUSE based filesystem:
 - `docker run --rm -it --cap-add SYS_ADMIN --device /dev/fuse sshfs`
- Give access to a single device:
 - `docker run -it --device=/dev/ttyUSB0 debian bash`

Give access to all devices:

- `docker run -it --privileged -v /dev/bus/usb:/dev/bus/usb debian bash`

1.6 Info

- `docker ps` shows running containers.
- `docker logs` gets logs from container. (You can use a custom log driver, but logs is only available for json-file and journald in 1.10).
- `docker inspect` looks at all the info on a container (including IP address).
- `docker events` gets events from container.
- `docker port` shows public facing port of container.
- `docker top` shows running processes in container.
- `docker stats` shows containers' resource usage statistics.
- `docker diff` shows changed files in the container's FS.
- `docker ps -a` shows running and stopped containers.

1.7 Import / Export

- `docker cp` copies files or folders between a container and the local filesystem.
- `docker export` turns container filesystem into tarball archive stream to STDOUT.

1.8 Executing Commands

`docker exec` to execute a command in container.

2. Images

A template or blueprint for docker containers.

2.1 Lifecycle

- `docker images` shows all images.
- `docker import` creates an image from a tarball.
- `docker build` creates image from Dockerfile.
- `docker commit` creates image from a container, pausing it temporarily if it is running.
- `docker rmi` removes an image.
- `docker load` loads an image from a tar archive as STDIN, including images and tags (as of 0.7).
- `docker save` saves an image to a tar archive stream to STDOUT with all parent layers, tags & versions (as of 0.7).

2.2. Info

- `docker history` shows history of image.
- `docker tag` tags an image to a name (local or registry).

2.3. Cleaning up

- `docker rmi` remove specific images.
- `docker gc` a tool to clean up images that are no longer used by any containers in a safe manner.

2.4. Load/Save image

- `docker load < my_image.tar.gz` load an image from file
- `docker save my_image:my_tag | gzip > my_image.tar.gz` save an existing image

2.5. Import/Export container

- `cat my_container.tar.gz | docker import - my_image:my_tag` import a container as an image from file
- `docker export my_container | gzip > my_container.tar.gz` export an existing container

3. Networks

A small def goes here

3.1. Lifecycle

- `docker network create`
- `docker network rm`

3.2. Info

- `docker network ls`
- `docker network inspect`

3.3. Connection

- `docker network connect`
- `docker network disconnect`

4. Registry & Repository

A repository is a hosted collection of tagged images that together create the file system for a container.

A registry is a host -- a server that stores repositories and provides an HTTP API for managing the uploading and downloading of repositories.

Docker.com hosts its own index to a central registry which contains a large number of repositories.

- `docker login` to login to a registry.
- `docker logout` to logout from a registry.
- `docker search` searches registry for image.
- `docker pull` pulls an image from registry to local machine.
- `docker push` pushes an image to the registry from local machine.

5. Volumes

Docker volumes are free-floating filesystems. They don't have to be connected to a particular container. You should use volumes mounted from data-only containers for portability.

5.1. Lifecycle

- `docker volume create`

- `docker volume rm`

5.2. Info

- `docker volume ls`
- `docker volume inspect`

6. Exposing ports

- `docker run -p 127.0.0.1:5HOSTPORT:SCONTAINER-PORT` --name CONTAINER -t docker_image mapping the container port to the host port using --
- `EXPOSE <CONTAINERPORT>` expose port CONTAINERPORT at runtime (see dockertile)
- `docker port CONTAINER SCONTAINERPORT` check the mapped port

7. Tips

7.1. Get IP address

- > `docker inspect some_docker_id | grep IPAddress`
| cut -d "" -f 4
or `install jq`
- > `docker inspect some_docker_id | jq -r '[0].NetworkSettings.IPAddress'`
or using a go template
- > `docker inspect -f '{{ (.NetworkSettings.IPAddress) }}' <container_name>`

7.2. Get port mapping

```
docker inspect -f '{{range $p, $conf := .NetworkSettings.Ports}}{{ $p }} -> {{index $conf 0}.HostPort}} {{end}}' <containername>
```

7.3. Find containers by regular expression

```
for i in $(docker ps -a | grep "REGEXP_PATTERN" | cut -f1 -d" "); do echo $i; done
```

7.4. Get Environment Settings

```
docker run --rm ubuntu env
```

7.5. Kill running containers

```
docker kill $(docker ps -q)
```

7.6. Delete old containers

```
docker ps -a | grep 'weeks ago' | awk '{print $1}' | xargs docker rm
```

7.7. Delete stopped containers

```
docker rm -v $(docker ps -a -q -f status=exited)
```

7.8. Delete dangling images

```
docker rmi $(docker images -q -f dangling=true)
```

7.9. Delete all images

```
docker rmi $(docker images -q)
```

7.10. Delete dangling volumes

```
docker volume rm $(docker volume ls -q -f dangling=true)
```

Exercises

Tiny project

Single-responsibility principle

Why we should decompose this project into tiny-tiny parts?

Wikipedia:

The single-responsibility principle (SRP) is a computer-programming principle that states that every class in a computer program should have responsibility over a single part of that program's functionality, which it should encapsulate. All of that module, class or function's services should be narrowly aligned with that responsibility [8].

Our Labs

Wine with X11

Raspbian on Docker

Security

References

References I



What is a Container? URL:

<https://www.docker.com/resources/what-container>.



packhelp. URL: <https://packhelp.com/plain-shipping-box/>.



64-bit memory layout. URL:

<https://www.berthon.eu/wiki/foss:wikishelf:linux:memory>.



32-bit memory layout. URL:

<https://unix.stackexchange.com/questions/31407/how-does-forking-affect-a-processs-memory-layout>.



Build a Docker Image just like how you would configure a VM. URL: <https://medium.com/platformer-blog/practical-guide-on-writing-a-dockerfile-for-your-application-89376f88b3b5>.

References II



List Of Docker Commands: Cheat Sheet. URL:
<https://phoenixnap.com/kb/list-of-docker-commands-cheat-sheet>.



Docker Commands - Complete Cheat Sheet. URL:
<https://linuxide.com/linux-how-to/docker-commands-cheat-sheet/>.



Single-responsibility principle. URL:
https://en.wikipedia.org/wiki/Single-responsibility_principle.