

The Container Security in Healthcare Data Exchange System

Bachelor's degree graduation project

Chih-Hsuan Yang

National Sun Yat-sen University

March 26, 2021
v1.0

Outline

- 1 Outcome
- 2 Related works
- 3 Current progress
- 4 Reference

Outcome

Outcome

No outcome.

Related works

Two papers

- A Measurement Study on Linux Container Security: Attacks and Countermeasures[1]
- Container-Based Cloud Platform for Mobile Computation Offloading[2]

Some Golang/Rust



- ① The next generation of C/C++
- ② High Concurrency, Memory Safe, Traits

Why?

High performance and secure server.

The docker-engine is written by Golang.

ASLR/KASLR/Finer-grained KASLR




```
Run /init.sh as init process
```

```
with arguments:
```

```
/init.sh
```

```
with environment:
```

```
HOME=/
```

```
TERM=linux
```

```
hostfs=./rootfs
```

```
mem=64M
```

```
kaslr: loading out-of-tree module taints kernel.
```

```
1694699525
```

```
random: fast init done
```

```
random: crng init done
```

```
→ 0326 git:(main) X less /proc/$$/maps
```

```
→ 0326 git:(main) X python -c 'print(1694699525)'
```

```
→ 0326 git:(main) X python
```

```
Python 3.8.5 (default, Jan 27 2021, 15:41:15)
```

```
[GCC 9.3.0] on linux
```

```
Type "help", "copyright", "credits" or "license" for more information.
```

```
>>> import math
```

```
>>> math.log(1694699525, 2)
```

```
30.658382356126655
```

```
>>> exit()
```

```
→ 0326 git:(main) X
```

Finer-grained KASLR


Finer-grained kernel address-space layout randomization[3]
Not merged on mainline yet.

New idea

Run container in UML?

<https://github.com/weber-software/diuid>

User-mode Linux



Capabilities

Table 3: Function of Security Mechanisms in Preventing Privilege Escalation Attacks

EDB-ID	CVE-ID	Security Mechanisms				
		Namespace	Cgroup	Capability	Secomp	MAC
Web App Layer						
43002	CVE-2017-15276			●		
40921	CVE-2016-9566			●		
42305	CVE-2017-6970			●		
40938	CVE-2014-6271			●		
Server Layer						
40768	CVE-2016-1247			●		
40678	CVE-2016-6663			●		
40450	CVE-2016-1240			●		
Kernel Layer						
41994	CVE-2017-7388					
43127 [*]						
43029 [*]	CVE-2017-5123					
40871	CVE-2016-8655					
40489						
40435						
44300						
40049	CVE-2016-4997			● NET_ADMIN ¹		
41458	CVE-2017-6074			● NET_ADMIN ¹		
43418	CVE-2017-1000112			● NET_ADMIN ¹		
41995	CVE-2016-9793			● NET_ADMIN ¹		
42887	CVE-2017-1000253			●		
42274	CVE-2017-1000566					
42275	CVE-2017-1000371			●		
42276	CVE-2017-1000379					
	CVE-2017-1000370					
40003					●	
39277	CVE-2016-0728					
39992	CVE-2016-1583			●	●	●
41762	CVE-2017-1575			●	●	●
41763	CVE-2017-1576			●	●	●
39166				●	●	●
39230	CVE-2015-8660			●	●	●
40847						
40616		●		●		
40611	CVE-2016-5195					
40839		●		●	●	
40838		●		●	●	
40759						
39772	CVE-2016-4357	●		●	●	
41999	CVE-2016-2384	●	●			

^{*} Security mechanism blocks the exploit.

^{*} Exploit bypasses all 5 security mechanisms.

¹ Exploit can achieve privilege escalation when the "NET_ADMIN" capability is included in the *cap_bset* of the caller process. Other exploits marked "•" in "Capability" column can only be successful when all 38 capabilities are included in the *cap_bset*. The "cap_bset" defines the highest privilege a process could reach.

Tim Hsu

- ① 「你研究了 capabilities 了嗎？」
 - 有試過，不能算研究。
- ② 「如果有 kernel exploit 能打穿 container 那你有那些 anti-exploit 的方式？」
 - Not to kernel: KASLR, SECCOMP, capabilities
 - Landed kernel: Encrypt container database, alert
 - UML? Hypervisor.

Tim Hsu

- ① 「先從只用現有機制 user land 的方式再往 kernel land 的方向」
- ② 「擋掉 kernel exploit 『可以從無法執行』或『可執行但不會成功』或『可執行但會被限制』等想法」
 - Capabilities, SECCOMP, AppArmor(Docker only)
 - No idea
 - Network Configuration, LSM


Interact with FHIR in docker

```
→ ~ curl -k -i -u 'fhiruser:change-password' 'https://localhost:9443/fhir-server/api/v4/$healthcheck'
HTTP/2 200
content-type: application/fhir+json
date: Thu, 25 Mar 2021 11:22:18 GMT
content-language: en-US
content-length: 123






{"resourceType":"OperationOutcome","issue":[{"severity":"information","code":"informational","details":{"text":"All OK"}}]}%
→ ~
```

Current progress

Map-reduce



Map-reduce

- Study:  $\frac{10}{\infty}$
- Container Security:  read some
- FHIR system:  Configured, can run.
- Combination:  Of course: 0
- 專題競賽暨成果展:  $\frac{48}{265} \sim 0.181$

Reference

References I

- [1] Xin Lin et al. “A Measurement Study on Linux Container Security: Attacks and Countermeasures”. In: ACSAC '18. San Juan, PR, USA: Association for Computing Machinery, 2018, 418–429. ISBN: 9781450365697. DOI: 10.1145/3274694.3274720. URL: <https://doi.org/10.1145/3274694.3274720>.
- [2] S. Wu et al. “Container-Based Cloud Platform for Mobile Computation Offloading”. In: *2017 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*. 2017, pp. 123–132. DOI: 10.1109/IPDPS.2017.47.
- [3] Jake Edge. 2020. URL: <https://lwn.net/Articles/812438/>.