

# The Container Security in Healthcare Data Exchange System

Bachelor's degree graduation project

Chih-Hsuan Yang

National Sun Yat-sen University

Advisor: Chun-I Fan

July 30, 2021

# Outline

- 1 Divide and Conquer
- 2 Why normal samples
- 3 How to find the fine-grained system calls?
- 4 Plans

# Flow chart



# Divide and Conquer

# 3 steps

- ① Sign and check
- ② Collect "Normal" samples
- ③ Enforce policy

# Is it can work?

- ✓ Sign and check
  - [docker scan](#)
  - [docker trust](#)
- ? Collect "Normal" samples
- ✓ Enforce policy in seccomp
  - Wrapper for [seccomp](#).

# seccomp - Secure Computing

Not use the LSM.

The seccomp and system call limitation are coupled.

## SECCOMP\_SET\_MODE\_FILTER

The system calls allowed are defined by a pointer to a Berkeley Packet Filter (BPF) passed via *args*. This argument is a pointer to a *struct sock\_fprog*; it can be designed to filter arbitrary system calls and system call arguments. If the filter is invalid, **seccomp()** fails, returning **EINVAL** in *errno*.


If **fork(2)** or **clone(2)** is allowed by the filter, any child processes will be **constrained** to the same system call filters as the parent. If **execve(2)** is allowed, the existing filters will be preserved across a call to **execve(2)**.

<https://man7.org/linux/man-pages/man2/seccomp.2.html>

## Why normal samples



# Microservices - Riemann sum



[https://commons.wikimedia.org/wiki/File:Riemann\\_sum\\_convergence.png](https://commons.wikimedia.org/wiki/File:Riemann_sum_convergence.png)

How to find the fine-grained system calls?

# LSMs

Might change the kernel. **Not inherit**  
Profiles and how to write it.

- AppArmor
- SELinux
- TOMOYO Linux

# LSMs

- AppArmor
  - Profiles
  - More fine-grained control in files.
- SELinux
  - Targeted, Role-Based-Access-Control, Multi-Layer/-Class Security
  - More fine-grained control in roles.
- TOMOYO Linux

# CGroups

It seems okay, but might be a little bit tricky to do it.

Reading, contributing.

Will be published in this weekend  
AT COSCUP.

## *Linux Kernel Scheduler Internals*

I am not a visionary. I'm an engineer. I'm happy with the people who are wandering around looking at the stars but I am looking at the ground and I want to fix the pothole before I fall in.

*TED 2018  
Linus Torvalds*

Ching-Chun (Jim) Huang and  
contributors

July 17, 2021



# CGroups

There are some people discuss when CGroups named CGroups from 'process containers'

- ① 2008 [cgroups: implement device whitelist cgroupslsm](#)
- ② 2011 [cgroup: syscalls limiting subsystem](#)
- ③ 2011 [Limiting system calls via control groups?](#)

*alonz: I wonder if it wouldn't be better to start from the other end of the solution space—small, incremental extensions to seccomp.*

# seccomp

- Advantage
  - Inherit
  - Block by each system call
- Disadvantage
  - Not cross platforms (CPU architecture)
  - Custom kernel

# Papers

## Intrusion Detection System for Applications Using Linux Containers [1]

- Toward Smart Moving Target Defense for Linux Container Resiliency [2]
- Smart Moving Target Defense for Linux Container Resiliency [3]
- Improving the Security of Microservice Systems by Detecting and Tolerating Intrusions [4]
- Malchain: Virtual Application Behaviour Profiling by Aggregated Microservice Data Exchange Graph [5]



# BPF/eBPF

Tracee: Tracing Containers with eBPF

DockerCon19 - eBPF Superpowers

But they are not using in **runtime** container security.

# Plans

# Plans in next week

- Use the BPF feature and fuzzing technique to collect the "normal" system calls in healthcare data exchange system.
- Keep survey those papers.

# References

- [1] Amr S. Abed, Charles Clancy, and David S. Levy. "Intrusion Detection System for Applications Using Linux Containers". In: *Security and Trust Management*. Ed. by Sara Foresti. Cham: Springer International Publishing, 2015, pp. 123–135. ISBN: 978-3-319-24858-5.
- [2] Mohamed Azab et al. "Toward Smart Moving Target Defense for Linux Container Resiliency". In: *2016 IEEE 41st Conference on Local Computer Networks (LCN)*. 2016, pp. 619–622. DOI: [10.1109/LCN.2016.106](https://doi.org/10.1109/LCN.2016.106).
- [3] Mohamed Azab et al. "Smart Moving Target Defense for Linux Container Resiliency". In: *2016 IEEE 2nd International Conference on Collaboration and Internet Computing (CIC)*. 2016, pp. 122–130. DOI: [10.1109/CIC.2016.028](https://doi.org/10.1109/CIC.2016.028).
- [4] José Flora. "Improving the Security of Microservice Systems by Detecting and Tolerating Intrusions". In: *2020 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*. 2020, pp. 131–134. DOI: [10.1109/ISSREW51248.2020.00051](https://doi.org/10.1109/ISSREW51248.2020.00051).
- [5] Mohammad Mahdi Ghorbani et al. "Malchain: Virtual Application Behaviour Profiling by Aggregated Microservice Data Exchange Graph". In: *2020 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*. 2020, pp. 41–48. DOI: [10.1109/CloudCom49646.2020.00004](https://doi.org/10.1109/CloudCom49646.2020.00004).